http://www.diva-portal.org

This is the published version of a paper published in .

Citation for the original published paper (version of record):

Sturm, B., Kanhov, E., Holzapfel, A., Hodgson, T., Huang, R S. (2026)
Introducing AI Music Studies
*Music and Data*, 1(1): 4-8
https://doi.org/10.1525/mad.2026.1.1.4

Access to the published version may require subscription.

N.B. When citing this work, cite the original published paper.

Permanent link to this version:
http://urn.kb.se/resolve?urn=urn:nbn:se:kth:diva-391900

BOB L. T. STURM, ELIN KANHOV, ANDRE HOLZAPFEL, THOMAS HODGSON, AND RUJING STACY HUANG

# Introducing AI Music Studies

Downloaded from online.ucpress.edu/mad/article-pdf/1/1/4/952294/mad.2026.1.1.4.pdf by guest on 03 November 2026

On 10–12 December 2024, The First International Conference in AI Music Studies (AIMS) convened in central Stockholm, Sweden, gathering sixty scholars from across Europe, Asia, Australasia, and the Americas, and from a deliberately wide spread of disciplines: music, social science, computer science, law, economics and history.[1] Conversations that animated the conference stretch back further still, to early 2023, when the rapid emergence of generative AI technologies began to outpace the scholarly vocabulary able to discuss and critically evaluate them. Around that time, English scholar Matthew Kirschenbaum, writing in *The Atlantic*, observed of his experience with ChatGPT: "We may quickly find ourselves facing a textpocalypse, where machine-written language becomes the norm and human-written prose the exception."[2] That observation came about a year after the CEO of online AI music service Boomy mused, "What does a hundred billion songs per year look like, and how does that shift the market? Whoever is there first and whoever is doing that – the spoils [are] going to be ridiculous."[3] Taken together, such pronouncements registered a sense—felt across multiple research communities by early 2023—that an appropriate scholarly response was no longer obvious, and that the disciplinary tools available for thinking about AI music in particular were poorly equipped. The first AIMS conference, in 2024, convened to address that gap and was followed, in 2026, by a second gathering at the University of Nottingham.[4] In the three years between those first conversations and this issue, a research community began to consolidate around the term *AI Music Studies*. This inaugural issue of *Music and Data* further advances this movement.

The seven articles of this issue were among those presented at the Stockholm conference. They take up, in different registers, the question that opened the proceedings: do AI-generated music and the emerging practices around it demand the founding of a new field, or are existing disciplines and their respective methodologies adequate? This question was not settled during the conference, and we do not propose to settle it in this issue. What can be said is that the case for dedicated, sustained scholarly attention to AI music has

*Music and Data*, Volume 1. © 2026 by The Regents of the University of California. electronic ISSN: 3070-765X, pp. 4–8. All rights reserved. Direct requests for permission to photocopy or reproduce article content through the University of California Press's Reprints and Permissions web page, https://online.ucpress.edu/journals/pages/reprintspermissions. DOI: https://doi.org/10.1525/mad.2026.1.1.4

only grown stronger. In the time since the first conference, such music has moved from a fringe novelty to a quotidian feature of the music industry: the first AI-generated tracks have appeared on (and have in some cases been removed from) national music charts;[5] major streaming services are reporting being inundated by submissions of what they claim is AI-generated music;[6] commercial detection services have emerged to identify AI music after the fact;[7] attribution services have arisen alongside them, promising to quantify which items of training data are responsible, and in what proportion, for any given generated output;[8] and a backlash, public and legal, is gathering pace.[9] Each of these developments is, in its own way, a reaction to the same underlying condition: a technology has emerged into the public sphere faster than any of our institutions—scholarly, legal, industrial—have been able to respond.

The difficulty of scholarship under such conditions is, at least, twofold. The first is that technological development almost always precedes critical reflection, and changes have already been largely absorbed in society by the time critique appears in print. The second, less obvious but perhaps more consequential, is that more immediate term responses among users of these platforms (*Is this really music? Is this really art?*) mask deeper anxieties that may prove more enduring over the *longue durée*: questions about labor, ownership, infrastructure, taste, aesthetics, and value, the political economy of training data, and how music is learned and taught in diverse cultural contexts. The work of AI Music Studies, as we envision it, is to confront these challenges, both keeping pace with technological change while also taking into account the effect it has on diverse musical worlds.

This work is no easy task. The study of AI music sits awkwardly across at least four kinds of analytical objects simultaneously: a body of technical artifacts (models, datasets, training pipelines, platforms) constantly in danger of becoming obsolete or disappearing; an industry with its own economic logic; a cultural forum with listeners and practitioners; and a discursive phenomenon shaped in roughly equal measure by hyperbole and cynicism. It is challenging for any single disciplinary tradition to address all four with equal attention and yet treating any one of them in isolation risks missing the way they in fact constitute one another. The technical decisions encoded in a generative model, for example, are not separable from the economic incentives that shaped its development; the discourse around AI music is not separable from the cultural anxieties it produces; and the aesthetic experience of an AI-generated track is not separable from what listeners know, or think they know, about how it was made. Furthermore, as AI music enters the public domain it becomes increasingly evident which unspoken values and ideas we have about music *not* generated with AI in terms of authenticity, aesthetics, provenance and value. Historically, new technological inventions and interventions have persistently challenged common perceptions of music and culture. What is perhaps different now with AI is the scale and pace of its development and its application in so many domains and contexts. What this issue offers, therefore, is not a unified method for navigating this terrain, but a working demonstration of what serious engagement across its registers can look like.

The seven articles that follow are methodologically heterogeneous, drawing on musicology, music theory, music information retrieval, computational creativity, science and technology studies, and the political economy of platforms. Their concerns range from the



Downloaded from online.ucpress.edu/mad/article-pdf/1/1/4/952294/mad.2026.1.1.4.pdf by guest on 03 November 2026

aesthetics of music's encounter with its own machines, through the rhetoric and practice of commercial AI music platforms, to the technical representations and datasets on which those platforms rest, and to the legal and economic architectures that increasingly govern them.

The issue opens with Alexander Rehding's "J.S. Bach and Artificial Intelligence: The View from Music Theory," which traces three moments when Bach has stood in for the project of music AI in the public imagination: Douglas Hofstadter's *Gödel, Escher, Bach* (1979), David Cope's 1997 "Bach vs. EMI" compositional contest, and the 2019 Google chorale-harmonizing Doodle powered by the Coconet machine learning model. Each, Rehding shows, exposes a mismatch between music theory's commitment to structural and architectural listening and the sound-to-sound, chord-to-chord listening that contemporary AI music tends to optimize for. The "rookie mistakes" that music theorists are quick to flag in AI-generated chorales—parallel fifths, off-tonic openings, missing tonal unity—turn out to be less errors than evidence of a different aesthetic project, with its own listening habits. In the process, Rehding makes the further claim that the Bach chorales as taught in seminar rooms are themselves a sanitized version of Bach's actual practice. What AI Bach exposes, on this reading, is as much about the unstated values of music theory itself as about the limits of machine learning.

From this vantage point on music's longer history with its machines, the issue turns to the rhetoric surrounding AI music in the present. Liam Pram and Fabio Morreale's "Opening Musical Creativity? Embedded Ideologies in Generative-AI Music Systems" draws upon material from platform creators and users (including themselves) to unpack the rhetoric that frames four specific commercial platforms as a "democratization" of music, an "unlocking of musical creativity," or even an inevitable technological development. Here we see how the threat to musicians or music artistry is overlooked in favor of how AI "empowers" or becomes a partner in helping one discover their untapped inner creative potential.

What the user actually does with these platforms, in practice, is the subject of Elin Kanhov and Bob L. T. Sturm's "Individualization, Platformization and Claims of Democratization: Motivations, Processes and Experiences of People and their Practices of AI Music." Through a thematic analysis of semi-structured interviews of nine users, Kanhov and Sturm reveal a rich set of motivations for why people turn to AI music platforms, and the diverse ways they use the platforms. From these perspectives, it is impossible to cast aside this mode of music-making as illegitimate or even thoughtless. While certainly some uses of these platforms principally seem to be tapping into online music streaming revenue, there are people who benefit in other ways from their engagement with these systems. Ultimately, however, this way of creating music is closely tied to corporate interests and a wider network of music platformization.

Beneath the platforms and their users sit the technical representations on which contemporary music AI is built. Landon Morrison and Andrew McPherson's "On the Reductivism of Disentangled Representations in AI Music Applications" critiques the practice of reducing sound and audio features to "disentangled representations" in digital music technologies. They argue that such reductions flatten musical representations and critically analyze how AI music applications reinforce correlations between words—in cases of AI in



Downloaded from online.ucpress.edu/mad/article-pdf/1/1/4/952294/mad.2026.1.1.4.pdf by guest on 03 November 2026

text-prompt systems—and their auditory representations. Via the AI audio synthesis technologies Magenta DDSP and RAVE and by drawing on (post)humanist philosophy and critical theory, the authors turn their critique into analytical action.

The same critical attention can be brought to bear on the datasets that train these systems, and Michèle Duguay and Johanna Devaney do so in "Constructing the Male and Female Singing Voice: Gender Encoding in AudioSet, MusicCaps, and MusicLM." They put vocal gender essentialism at the core of this investigation, asking to what extent these datasets and models perpetuate a strict alignment between specific sonic features of the voice and a binary gender framework seen as equivalent to biological sex. Their results demonstrate how pitch and spectral features consistently align with gender labels, and that these differences increase even more when using a synthesis model compared to the datasets that it has been trained upon. The authors recommend changes in labeling strategies but also consider the possibility of refusal by affected groups to be properly represented by future surveillance and control technologies.

How value flows through AI music systems—and to whom—is the subject of Oliver Bown's "Training Data Attribution in Music AI: Analyzing Strategies and Their Discourses of Value." Bown looks at a growing market for solving a problem created by AI music generation systems: the attribution (and possibly remuneration) of rights holders of training data. Bown conceptualizes six "determinants of value" of music, and tests their alignment with different attribution services and claims made by their developers. Bown notes the incongruity of intellectual property law as protecting specific expressions, not style shared among many expressions, and yet attribution effectively rewarding stylistic similarity between model output and training data. Bown warns against the techno-solutionism of automated attribution, bringing more "opacity and the very prominent threat of being gamed, undermining rather than serving small independent creatives, underlying principles of fairness and public benefit."

The issue closes with Eric Drott's "The Privatization of AI Music Regulation and the Assetization of the Voice," which describes how decisions about copyright infringements are increasingly made by corporations behind closed doors in a process of privatization of copyright enforcement. Using the example of the human voice, he illustrates how such privatized systems and recent public legislation drafts converge into a reinforcement of the identification of human voice and self, by sharing the expressed needs for technical measures that codify, uphold, and potentially extend this identification. This convergence results in the voice itself becoming a novel economically exploitable asset, related to and protected by mechanisms of privatized copyright enforcement. It is a fitting close to the issue: a piece that reaches outward from AI music into the broader political economy of cultural regulation, and that sets up questions about ownership, governance, and the public interest that the next issues of this journal will have ample occasion to take up.

**Finally, a brief procedural note: each article in this issue underwent two rounds of** double-blind peer review, with revision and re-review between rounds, and we are indebted to the generous community of reviewers whose work is by convention uncredited but without whom none of what follows would have been possible. We are grateful, too, to the participants and sponsors of The First International Conference in AI Music Studies:



Downloaded from online.ucpress.edu/mad/article-pdf/1/1/4/952294/mad.2026.1.1.4.pdf by guest on 03 November 2026

Riksbankens Jubileumsfond, Kungliga Musikaliska Akademien, and the European Research Council (ERC-2019-COG No. 864189); and to the University of California Press for hosting *Music and Data* at what we believe is a genuinely opportune moment for the questions this issue takes up.

## NOTES

1. The conference website and collected materials can be found here: AI Music Studies 2024, "AIMS: The First International Conference in AI Music Studies," https://boblsturm.github.io/aimusicstudies2024/ (accessed May 27, 2026).

2. Matthew Kirschenbaum, "Prepare for the Textpocalypse," *The Atlantic*, March 8, 2023, https://www.theatlantic.com/technology/archive/2023/03/ai-chatgpt-writing-language-models/673318/.

3. Valerio Velardo, "In Conversation with Alex Mitchell - Boomy," YouTube video, 01:07:28, April 7, 2022, https://youtu.be/iyTJF7b6BwE (accessed June 3, 2026).

4. These conferences grew out of conversations within the ERC-funded MUSAiC project (Music at the Frontiers of Artificial Creativity and Criticism, ERC-2019-COG No. 864189) at KTH Royal Institute of Technology in early 2023, and a subsequent position paper drafted by members of the project and several collaborators. After feedback from established music informatics venues that—in retrospect—helped clarify the case for a dedicated forum, the paper was posted publicly and planning for the Stockholm conference began. See Bob L. T. Sturm et al., "AI Music Studies: Preparing for the Coming Flood," *Proceedings of AIMC 2024 (09/09–11/09)* (August 2024): https://aimc2024.pubpub.org/pub/ej9b5mv1. Project website: https://musaiclab.wordpress.com/. See also https://www.aims2026.com/ (accessed June 4, 2026).

5. See Stuart Hughes, "Song Banned from Swedish Charts for Being AI Creation," *BBC News*, January 15, 2026, https://www.bbc.com/news/articles/cp829jey9z70 (accessed June 2, 2026).

6. Kristin Robinson, "Deezer Says Number of Fully AI-Generated Songs Delivered Daily has Doubled since January," *Billboard*, April 16, 2025, https://www.billboard.com/pro/deezer-fully-ai-generated-songs-uploaded-daily/.

7. For instance, Jesper Wendel, "Deezer Confirms Demonetization of up to 85% of AI-Music Streams Due to Fraud and Moves to Sell AI-Detection Technology," *Deezer Newsroom*, January 29, 2026, https://newsroom-deezer.com/2026/01/ai-generated-music-deezer-selling-detection-tool/; Laura Cros Vila et al., "The AI Music Arms Race: On the Detection of AI-Generated Music," *Transactions of the International Society for Music Information Retrieval* 8, no. 1 (2025): 179–94.

8. See "STIM Launches the World's First AI Liscense for Music," *STIM*, September 8, 2025, https://www.stim.se/en/news/stim-launches-the-worlds-first-ai-license-for-music (accessed May 27, 2026).

9. Woodrow Hartzog and Jessica Silbey, "How AI Destroys Institutions," *UC Law Journal* 77 (2026): 727–70.

Downloaded from online.ucpress.edu/mad/article-pdf/1/1/4/952294/mad.2026.1.1.4.pdf by guest on 03 November 2026